



# Kubernetes Flow – A Deep Dive



## Kubernetes Architecture Overview

At a high level, Kubernetes has a **control plane** and a **set of worker nodes**:



### Control Plane Components

| Component          | Purpose                                  |
|--------------------|------------------------------------------|
| kube-apiserver     | Entry point for all Kubernetes comma     |
| etcd               | Key-value store for all cluster data (cc |
| kube-scheduler     | Assigns pods to nodes.                   |
| controller-manager | Manages controllers (replication, end    |



### Node Components

| Component         | Purpose                                |
|-------------------|----------------------------------------|
| kubelet           | Runs on each worker node, ensures co   |
| kube-proxy        | Maintains networking rules and forwar  |
| container runtime | (e.g., containerd or Docker) Runs cont |



## Lifecycle of a Pod Creation (End-to-End Flow)

Let’s see what happens when a user runs:

```
bash
```

```
kubectl apply -f pod.yaml
```



Copy



Edit

1

`kubectl` → `kube-apiserver`

- `kubectl` is the CLI tool.
- It serializes `pod.yaml` into a REST API request.
- It sends a POST request to the API server, e.g., `POST /api/v1/namespaces/default/pods`.

## 2 kube-apiserver → etcd

- The `kube-apiserver` is the **front controller** of the control plane.
- It **validates** the pod spec (authentication, authorization, admission control).
- If valid, it **writes the Pod object** into `etcd`.

`etcd` stores:

yaml



Copy



Edit

```
/api/v1/namespaces/default/pods/mypod
```

## 3 kube-scheduler Watches Unscheduled Pods

- The **scheduler** continuously **watches the API server** for Pods with `spec.nodeName: null`.
- It selects a node based on:
  - Resource availability
  - Affinity/anti-affinity
  - Taints/tolerations
  - Node selectors, etc.

It then **updates the Pod spec** with:

yaml



Copy



Edit

```
spec: nodeName: "worker-node-2"
```

This change is saved back to `etcd` via the API server.

## 4 kubelet Watches Assigned Pods

- Each node runs a **kubelet**, which watches the API server for **Pod objects scheduled to its node**.
- When it sees a new pod, the kubelet:
  - Pulls the container image
  - Creates container via **container runtime**
  - Mounts volumes, sets up networking
  - Starts the container(s)

## 5 Pod is Running 🎉

- The kubelet continuously:
  - Probes liveness/readiness
  - Sends status updates to API server

## 6 Controllers in Background

- **ReplicaSetController**: Watches Deployments and maintains number of replicas.
- **JobController, DaemonSetController, etc.**, all use the same pattern:
  - Watch → Reconcile → Update desired state.

## Detailed Example: What Happens Internally?

Imagine this Pod spec:

yaml



Copy



Edit

```
apiVersion: v1 kind: Pod metadata: name: hello-world spec: containers: - name: hello image: busybox command: ["echo", "Hello Kubernetes!"]
```

## Internals (Step-by-step)

1. User runs: `kubectl apply -f pod.yaml`
2. `kubectl` sends HTTP request to API server.
3. **API server** does:
  - Authentication (e.g., via TLS certs, tokens).
  - Authorization (e.g., RBAC).
  - Admission Control (validations, limits).
  - Writes to **etcd** (desired state).
4. **Scheduler** detects unscheduled Pod and assigns it to `node-x`.
5. **API server** updates etcd with assigned node.
6. **kubelet on node-x** detects new Pod spec via API.
7. **kubelet** starts the Pod using `containerd`.
8. Pod is created and started.
9. **kubelet** reports status to the API server, which stores it in etcd.
10. `kubectl get pod hello-world` shows:

sql



Copy



Edit

| NAME        | STATUS  | AGE |
|-------------|---------|-----|
| hello-world | Running | 10s |



# Kubernetes Flow Summary (Diagram)



## Additional Notes for Students

- `kubectl get pod -o yaml` lets you see how the Pod looks in etcd.
- Every component talks only to `kube-apiserver`.
- All reconciliation loops are controller-driven: "**desired state vs actual state.**"
- `kubelet` is responsible only for its node.
- `kube-proxy` handles **services and load-balancing** using iptables or IPVS.